


# Hadoop 2 and YARN



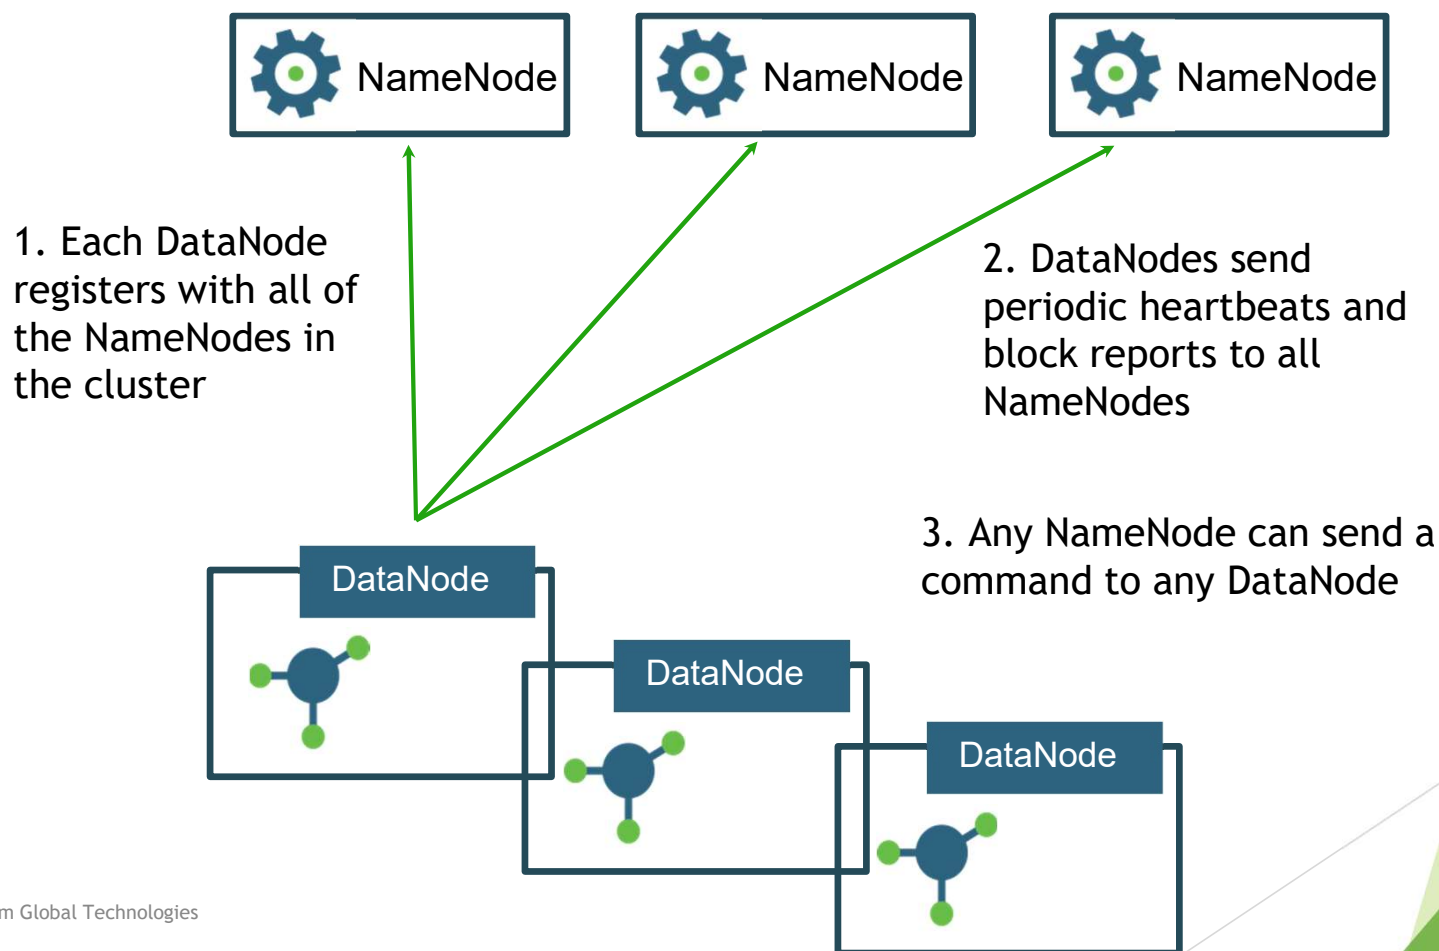
# Topics Covered

- About HDFS Federation
- Overview of HDFS High Availability
- Quorum Journal Manager
- Configuring Automatic Failover
- About YARN
- The Components of YARN
- Lifecycle of a YARN Application
- A Cluster View Example
- Introduction to Apache Slider
- *Lab: Running a YARN Application*

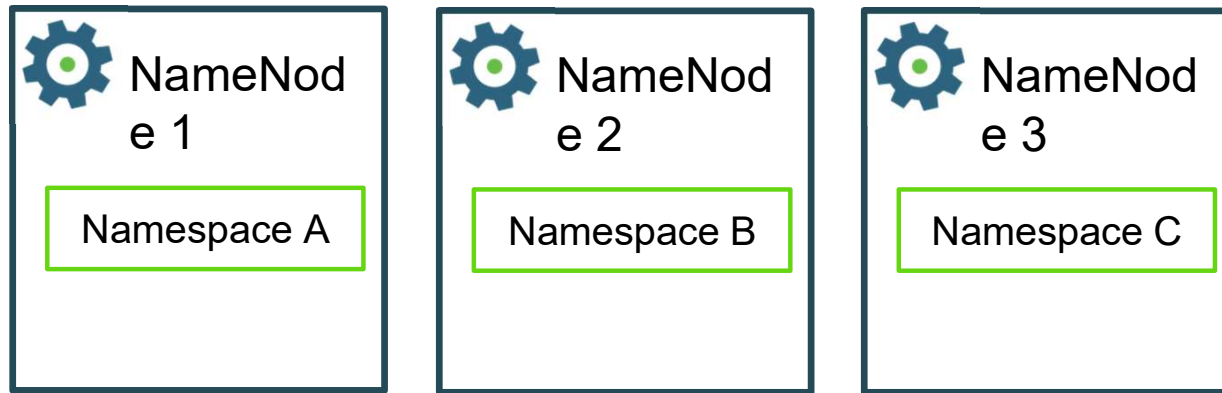
## About HDFS Federation

- According to Merriam-Webster's dictionary: a ***federation*** is an organization or group within which smaller divisions have some degree of internal autonomy
- **HDFS Federation** refers to the ability of NameNodes to work independently of each other

# Multiple Federated NameNodes



# Multiple Namespaces

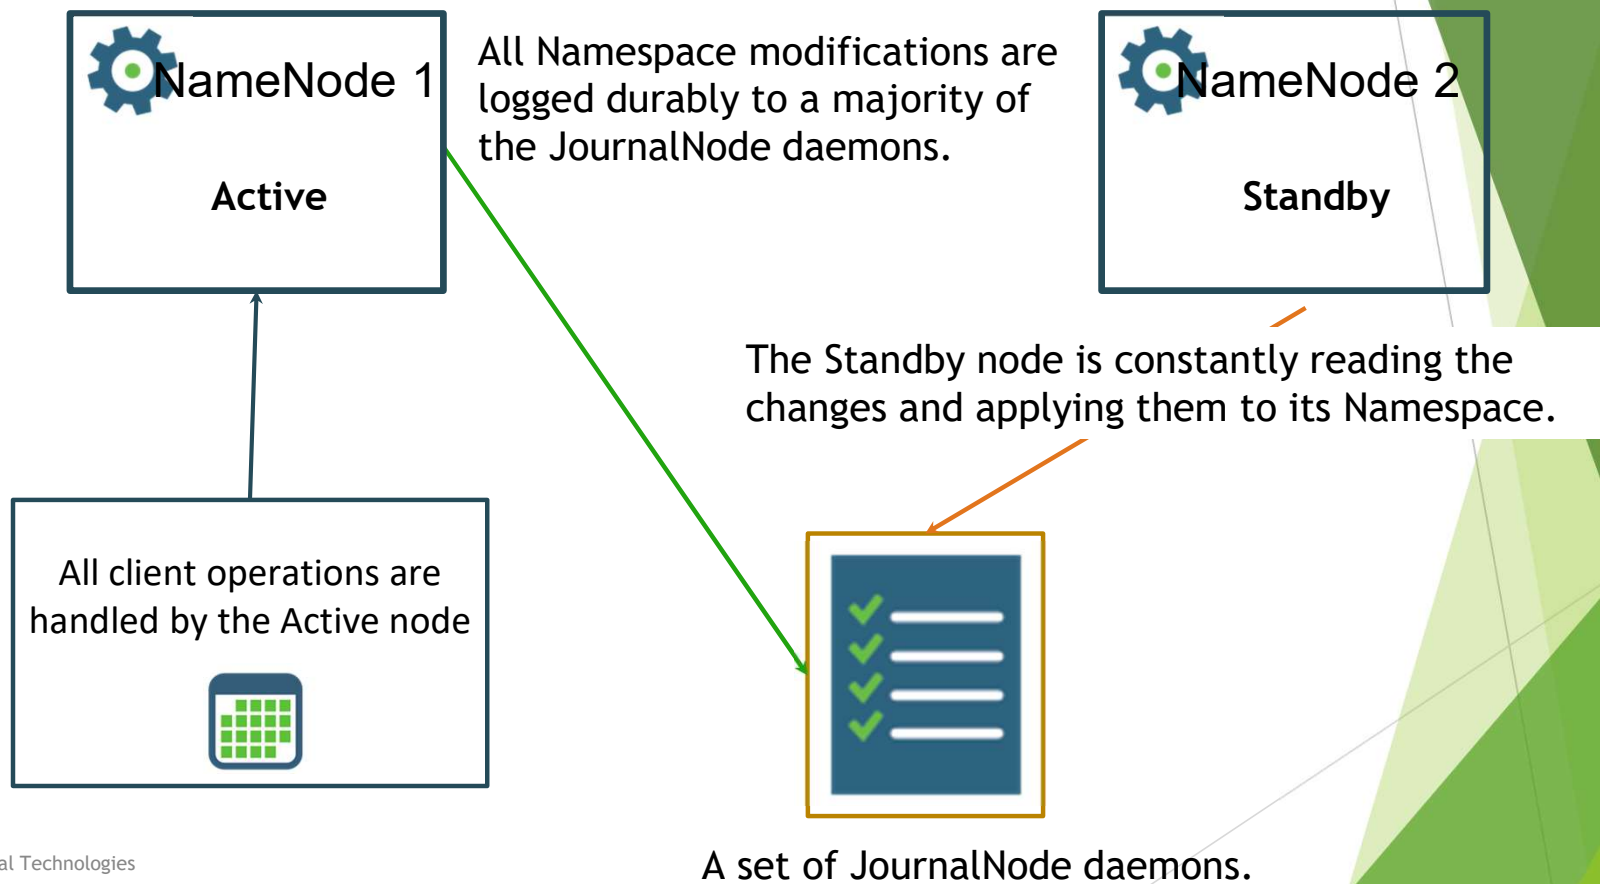


- Files and directories belong to a Namespace
- Prior versions of Hadoop only had a single Namespace
- Hadoop 2.x allows for multiple Namespaces
- A NameNode manages a single Namespace Volume


## Overview of HDFS HA

- In prior versions of Hadoop, the NameNode was a single point of failure that required additional tools to achieve HA
- In Hadoop 2.x, NameNode HA can be achieved using the built-in Quorum Journal Manager framework, Zookeeper, and the new Zookeeper Failover Controller processes

# Quorum Journal Manager



# Configuring Automatic Failover



ZKFC holds a lock to the Active NameNode. If that NameNode fails, the lock is made available.



The ZooKeeper daemons determine if a NameNode has failed. It also provides the lock that the ZKFC uses.



ZooKeeper instances.



## About YARN

YARN = Yet Another Resource Negotiator

YARN splits up the functionality of the JobTracker in Hadoop 1.x into two separate processes:


- **ResourceManager:** for allocating resources and scheduling applications
- **ApplicationMaster:** for executing applications and providing failover

## Open-source YARN Use Cases

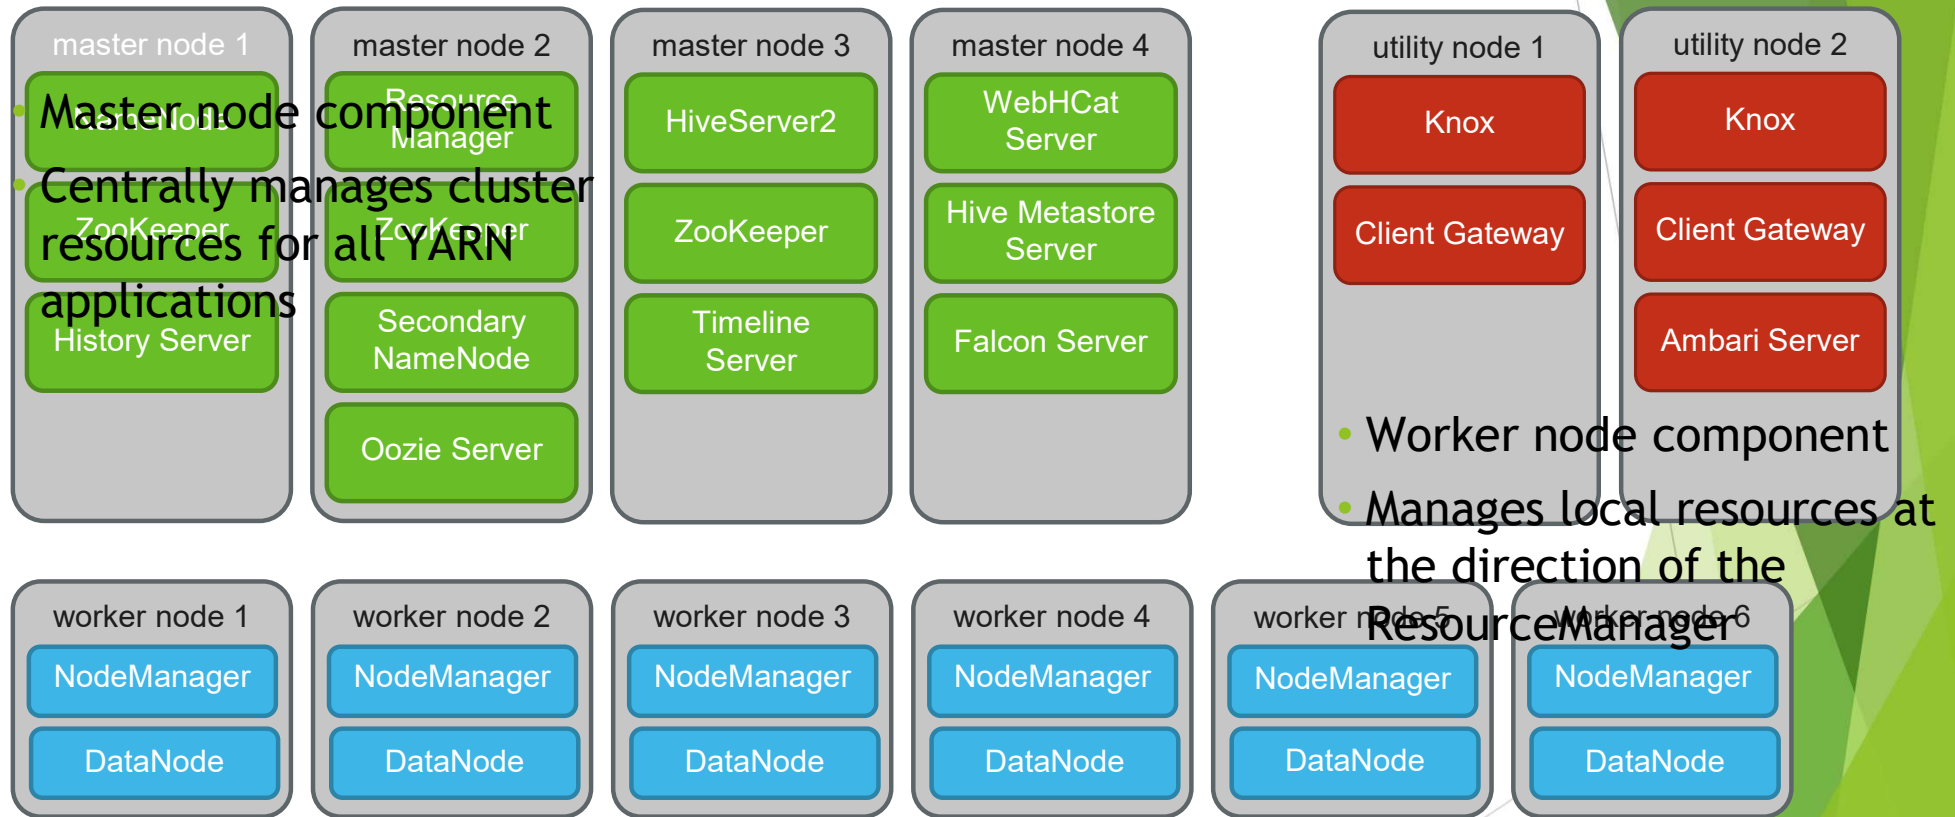
- **Tez:** improves the execution of MapReduce jobs
- **Slider:** deploy existing frameworks on YARN
- **Storm:** for real-time computing
- **Spark:** a MapReduce-like cluster computing framework designed for low-latency iterative jobs and interactive use from an interpreter
- **Apache Giraph:** a graph-processing platform

# The Components of YARN

The **ResourceManager** communicates with the **NodeManagers**, **ApplicationMasters**, and **Client** applications



# YARN Architecture - Big Picture View

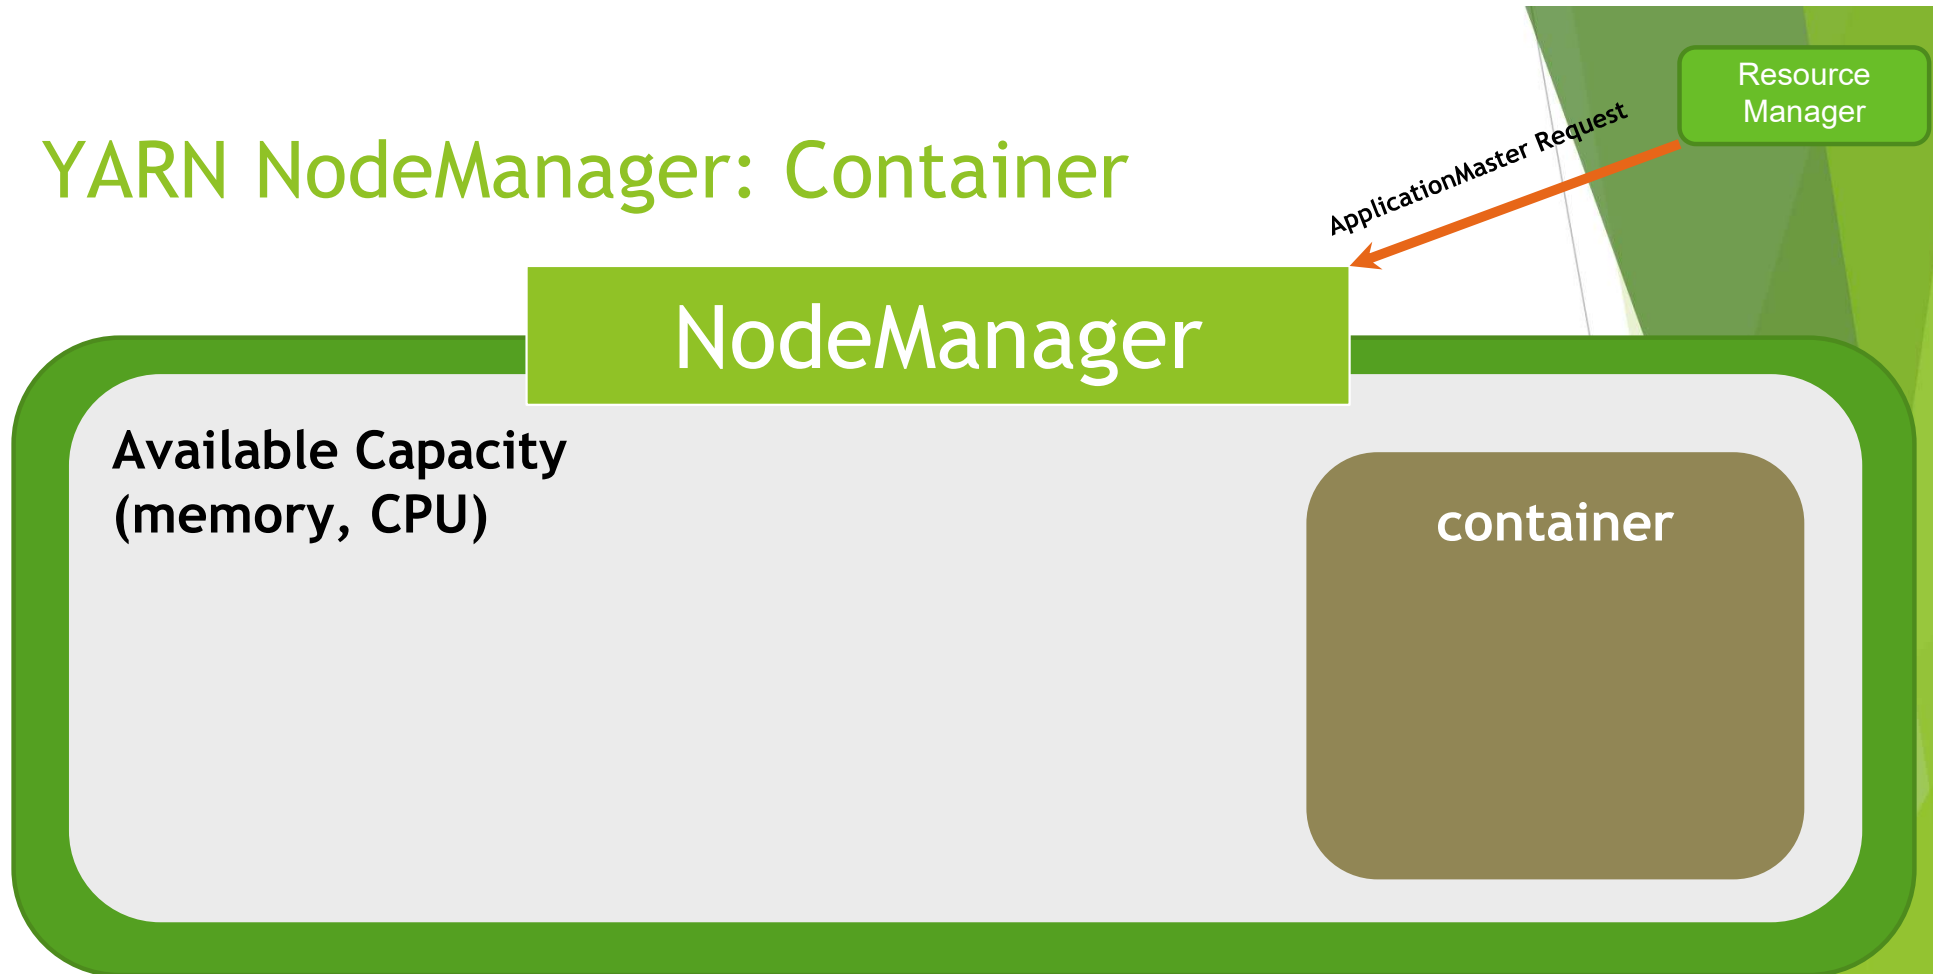


# YARN NodeManager

## NodeManager

- **Manages local CPU and RAM resources on behalf of requesting services**
- **Tracks node health and communicates status to the ResourceManager**

# YARN NodeManager: Container

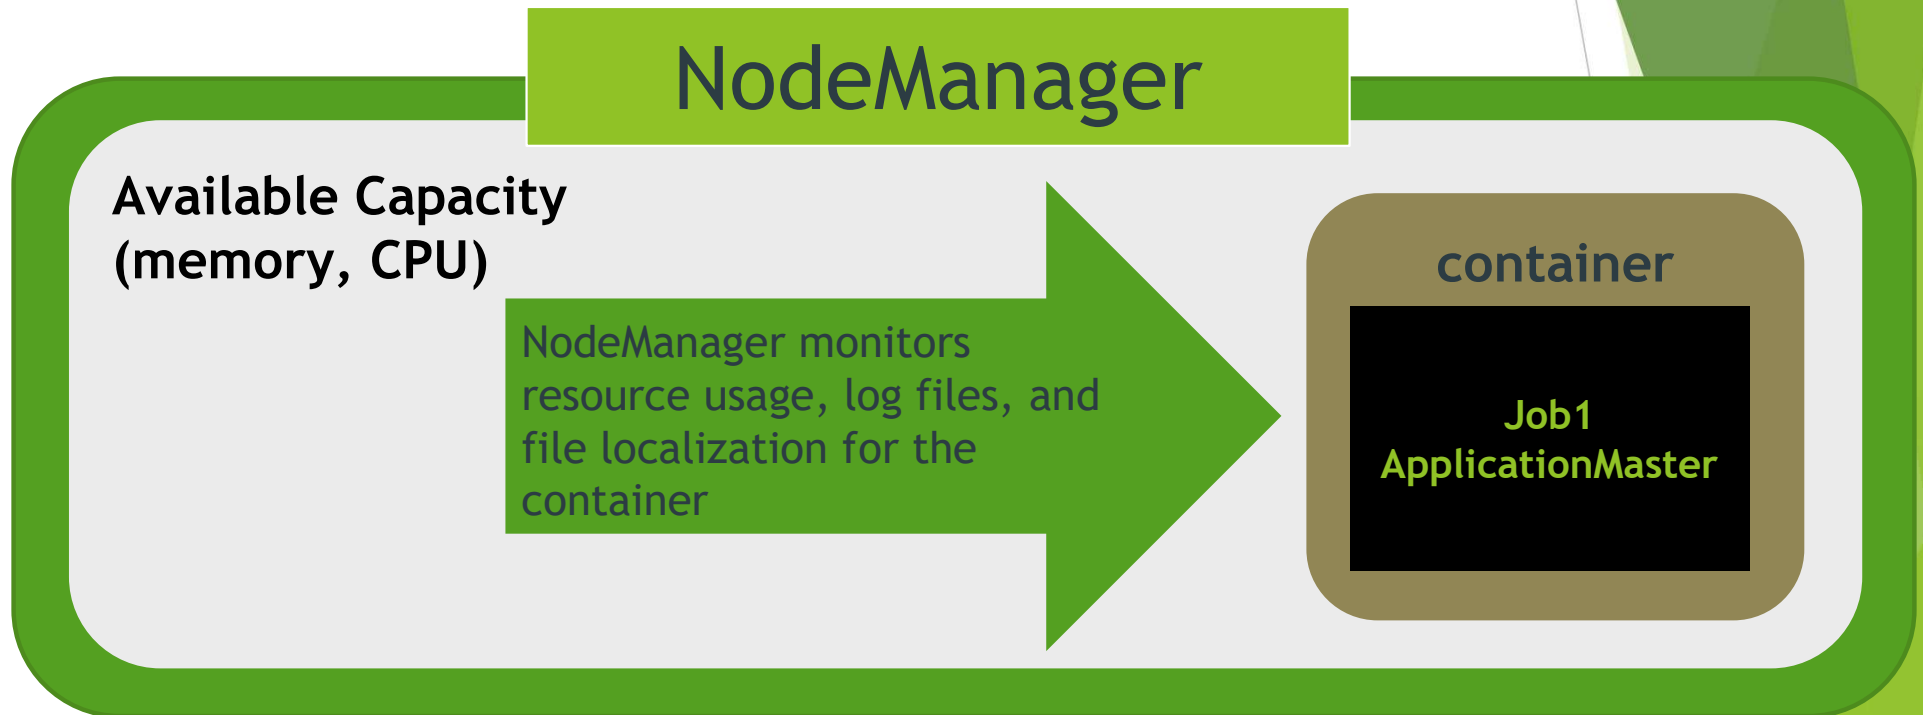


# Containers Defined

## container

- Allocated RAM and CPU cores by the NodeManager
- Runs ApplicationMaster job
- New container spawned for each discrete job task

# YARN NodeManager: ApplicationMaster



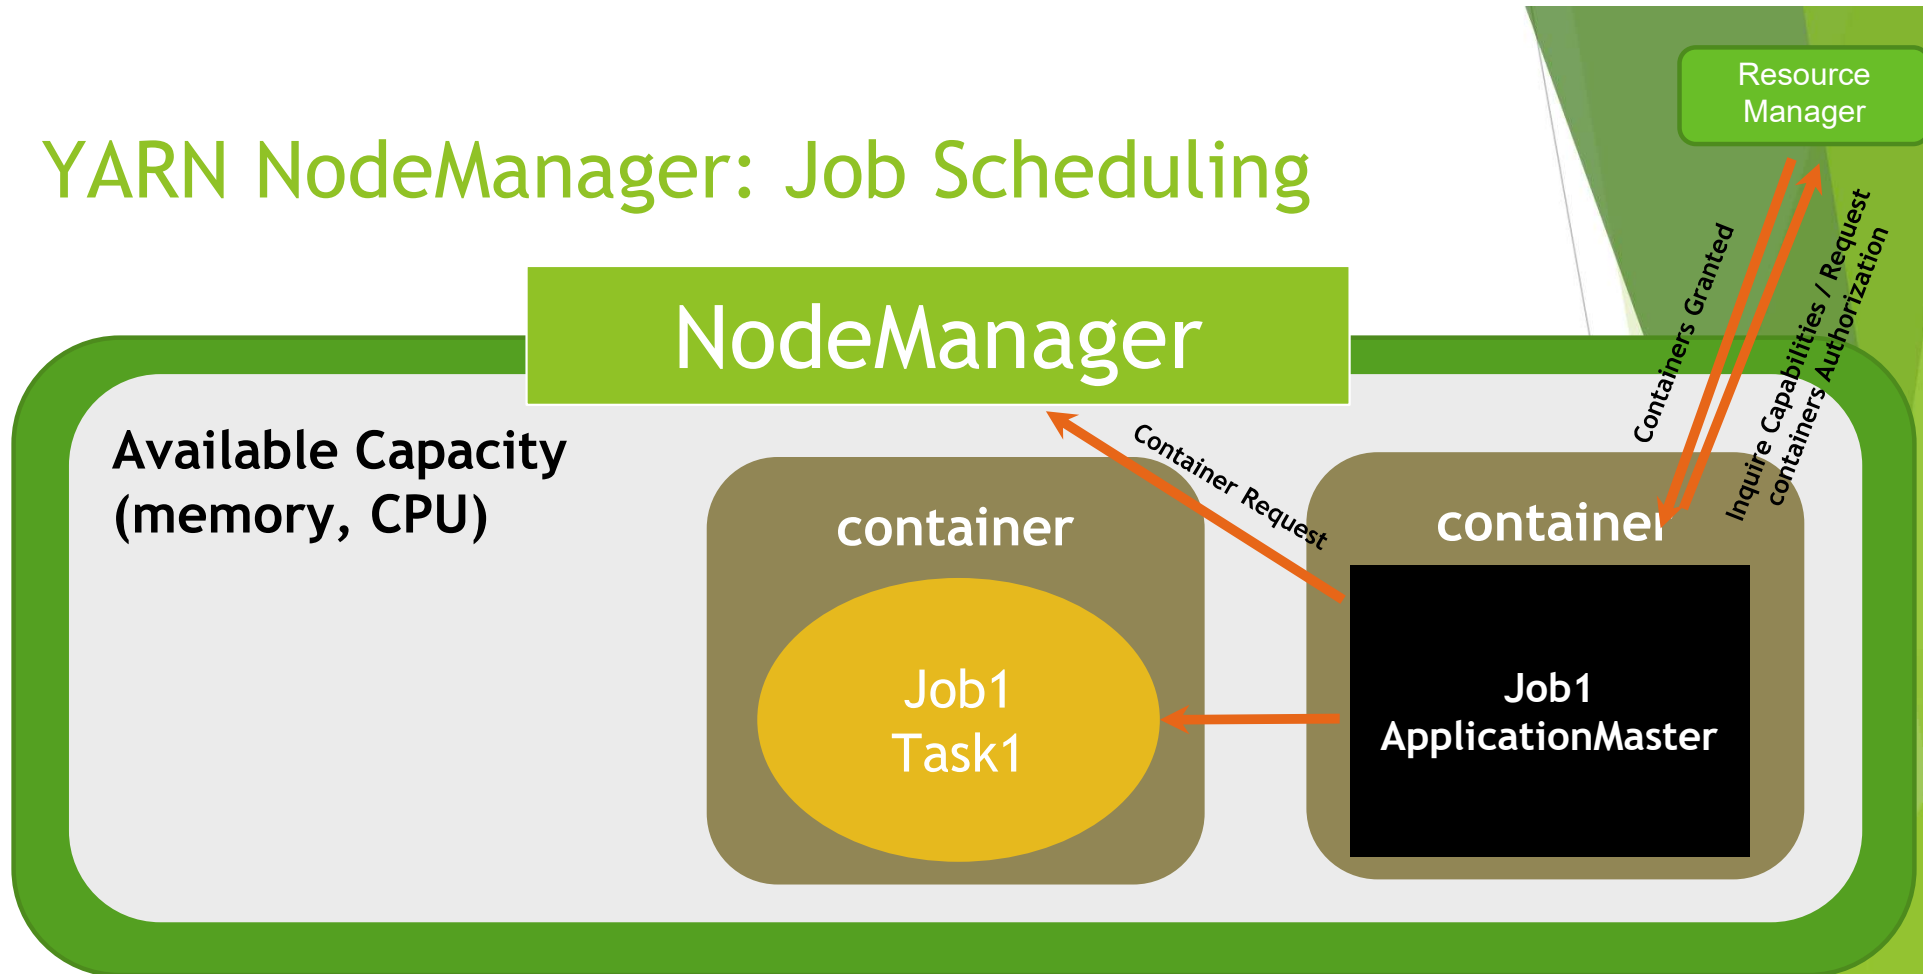


## ApplicationMasters Defined


### ApplicationMaster

- Bootstrap process for YARN applications
- Negotiates for resources with ResourceManager
- Works with NodeManagers to configure and execute containers and monitors application resource consumption
- Provides application fault tolerance and thus significant horizontal scale capabilities

# YARN NodeManager: Job Scheduling



## YARN NodeManager: Additional Tasks



# Example Multi-node Resource Allocation Scenario



Default behavior is to move processing to data rather than copy data to nodes with available processing capacity

# YARN ResourceManager (Master Node)

## ResourceManager

### Scheduler

- Controls global cluster resource usage
- Configurable by the Hadoop Administrator
- Enables multitenancy and Service Level Agreements

### Node Management

- Monitor NodeManager state
- Submit ApplicationMaster Requests
- Verify container launch
- Monitor ApplicationMaster state

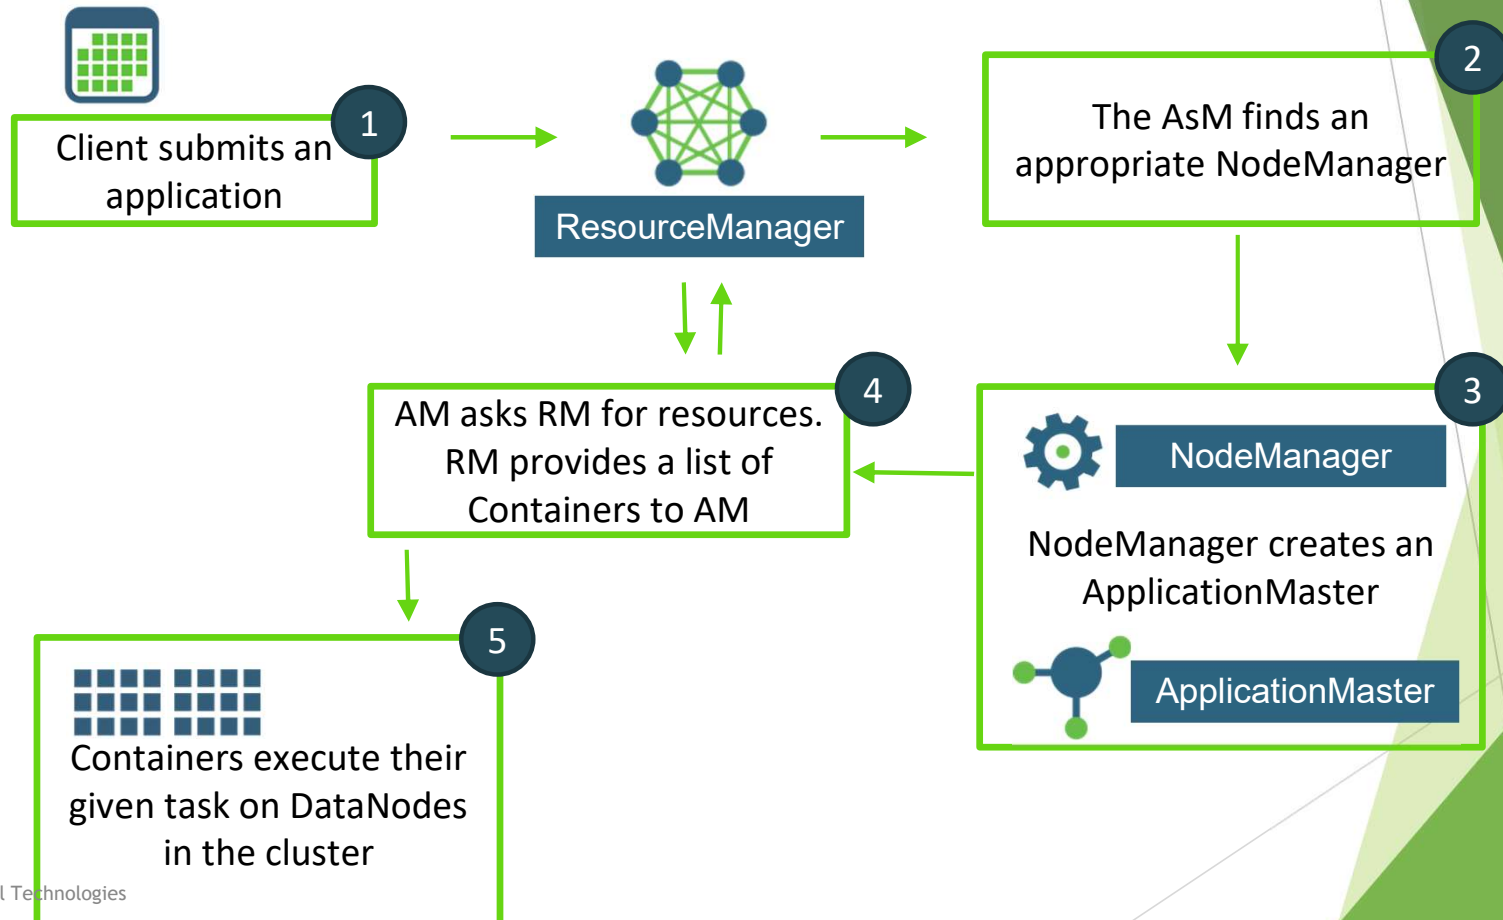
### Security

- Web Application Proxy
- User ACLs
- Manages tokens to ensure validity of all container requests made by ApplicationMasters

# YARN Component Summary

ResourceManager	NodeManager	Container	ApplicationMaster
Schedule global resources	Manage local memory and CPU allocation	Allocated RAM and CPU cores by NodeManager	YARN application bootstrap process
Enable multitenancy			Negotiate resources
Enable SLA enforcement			Provide application fault tolerance
Monitor and manage NodeManagers	Track and report on node health		Work with NodeManager for container restart
Monitor and manage ApplicationMasters	Manage file localization for containers	Run ApplicationMasters and job tasks	
Monitor containers globally	Monitor and manage local containers		Monitor job tasks and containers across cluster
Manage ACLs			
Manage Tokens			

# Lifecycle of a YARN Application



# A Cluster View Example

